

## Prospector

# An AI agent that finds your next *customers.*

Give it a goal in plain English. It plans, searches real maps and websites, recovers when a step fails, scores every lead with reasons, sets up a live AI receptionist you can call for the best ones, and learns what a good lead means to you.



# Every sales list today is either *slow or made up.*

## THE PROBLEM

1

### Slow

Someone searches Maps business by business, opens every website, and guesses who might buy.

2

### Made up

Ask a chatbot for 30 leads and it invents clinics, phone numbers and doctors that do not exist.

3

### Unconvincing

A spreadsheet row does not show a business owner what an AI receptionist would do for them.

## WHAT PROSPECTOR DID ON REAL RUNS IN PUNE

50

real dental clinics found in one run, each one traced back to OpenStreetMap


3

failed tool calls recovered from in the physiotherapy run, still finished in 18 of 20 steps

0

businesses or facts invented. Anything not observed is shown as "not observed"

# Plan, act, observe, re-plan: *a real agent loop.*



## TOOLS THE AGENT CAN CALL

geocode

search\_businesses

widen\_area

enrich

score

deep\_research

write\_hooks


prepare\_receptionists

finish

- **No agent framework.** No LangChain, no CrewAI. The loop, tool registry and rule validator are plain Python you can read.


- **Wrap-up guard.** If searches stop finding new businesses, it stops searching and always finishes the pipeline.

# Watch it think. Then *call* the lead's AI receptionist.



## Mission control

A live map fills with pins, the plan ticks off, and every thought, tool call and failure streams in as it happens.




## Live voice call with the Gemini Live API

The persona uses only facts the agent observed. Barge-in, live captions, takes messages, and replies in Hindi or Marathi. The browser gets a one-use token; the API key never leaves the server.

*Screenshots: the app running on the end-to-end test data set.*

# It learns what a good lead means to you, *and proves it.*



1

**You label leads Good or Bad**

One key per lead on the review page (G / B).

2

**The LLM proposes new scoring rules**

It reflects on the leads it scored wrong.

3

**Each rule is tested on unseen leads**

Rules are picked on one split and graded on a separate one.

4

**Kept only if accuracy goes up**

Otherwise rejected, and the reason is shown.

87

unit and  
API tests

7

end-to-end  
browser specs

Python

FastAPI

Next.js 16

Gemini

Gemini Live

OpenStreetMap

Screenshot: the app running on the end-to-end test data set.